

AI Targeting & Triage Agent

Agentic Version — Proof of Concept
Requirements Document

Field	Details
Document Type	POC Requirements Specification
Version	1.0
Status	Draft for Review
Prepared By	AI Architecture Team
Date	30 March 2026
Classification	Confidential – Internal Use Only

1. Overview

This document defines an Agentic AI Proof of Concept (POC) for a Targeting & Triage system. The system uses multiple specialised agents to deliver intelligent account prioritisation and sales enablement.

The system is designed to:

- Identify high-priority accounts
- Evaluate upsell potential
- Generate prioritised recommendations
- Provide explainable reasoning
- Highlight uncertainty and conflicts

2. Objective

Enable a goal-driven, multi-agent system that can autonomously support the sales process through intelligent data analysis, prioritisation, and decision support.

Core capabilities:

- Analyse account data from multiple sources

- Prioritise opportunities using ML and LLM signals
- Recommend targeted actions for sales teams
- Support human-in-the-loop decision-making

3. Agent Definitions

The system is composed of six specialised agents, each with a defined responsibility and data contract. The agents collaborate in an orchestrated pipeline to deliver end-to-end account prioritisation.

3.1 Planner Agent

Section	Details
Responsibility	• Interpret user request • Break task into sub-steps • Orchestrate downstream agents
Example Input	• User request: "Show top 30 accounts in NL & LU"
Example Output	• Fetch data from Data Agent • Score accounts via Scoring Agent • Prioritise with Reasoning Agent • Explain results via Formatting Agent

3.2 Data Agent

Section	Details
Responsibility	• Retrieve and prepare account data • Ensure data consistency and integrity
Input	• Filters: country, segment, date range
Output	• Structured, validated account dataset

3.3 Scoring Agent (ML)

Section	Details
Responsibility	• Generate buy/upsell propensity scores using ML models
Output Fields	• propensity_score: Numeric buy/upsell likelihood • confidence_level: Model confidence indicator

3.4 Reasoning Agent (LLM)

Section	Details
Responsibility	- Assign priority bucket (A / B / C) - Interpret contextual signals - Generate human-readable explanations for each recommendation

3.5 Validation Agent

Section	Details
Responsibility	- Detect conflicts between ML and LLM outputs - Evaluate consistency of recommendations
Conflict Rules	- ML High + LLM Low → Conflict flagged - ML Low + LLM High → Conflict flagged - Conflicts surfaced to output; not auto-resolved

3.6 Formatting Agent

Section	Details
Responsibility	- Produce structured, consistent output payload - Ensure all required fields are present and valid - Format for API response and downstream consumption

4. Data Requirements

The system ingests data from the following source files. All data must maintain referential integrity and valid cross-file relationships.

Source File	Description
accounts.csv	Core account master data
opportunities.csv	Sales opportunity records
snowflake_metrics.csv	Product usage and engagement metrics
external_funds.csv	External funding signals
conferences.csv	Conference and event engagement data

Data quality requirements:

- Referential integrity must be maintained across all source files
- Valid relationships between accounts, opportunities, and metrics
- No orphaned records or broken foreign key references

5. Feature Requirements

The system must compute and expose the following derived features as inputs to the Scoring and Reasoning Agents:

Feature Category	Feature Name	Description
Sales Performance	win_rate	Historical deal win percentage
Sales Performance	avg_deal_size	Average closed deal value
Product Coverage	product_count	Number of active products
Revenue	revenue_concentration	Revenue distribution across accounts
Growth	growth_metrics	Period-over-period revenue growth indicators
Service	service_penetration	Service adoption depth across account
Activity	launch_indicators	New product / campaign launch signals
Engagement	engagement_score	Composite account engagement index

6. Execution Flow

The agents execute in the following sequence for each scoring request:

1. Planner Agent receives and interprets the user request
2. Data Agent retrieves and validates account data
3. Scoring Agent computes ML-based propensity scores
4. Reasoning Agent assigns priority bucket and generates rationale
5. Validation Agent checks for ML vs LLM conflicts
6. Formatting Agent prepares the structured output payload
7. Response returned to the calling system or user

7. Output Requirements

Each account in the output must include the following fields:

Field	Type	Description
account_id	String	Unique account identifier
priority_bucket	Enum (A/B/C)	LLM-assigned priority tier
ml_score	Float	ML propensity score
confidence_level	Float	Model confidence indicator
conflict_flag	Boolean	ML vs LLM conflict indicator
status	String	Current account status
rationale_text	String	Human-readable explanation of prioritisation

8. API Requirements

POST /score_accounts

GET /account_view/{id}

Returns a detailed view of a single account including full scoring context.

Direction	Field	Description
Input	id (path param)	Unique account identifier
Output	account_detail	Full account view with scores, rationale, and conflict flags

9. Decision Rules

The following rules govern how ML and LLM outputs are combined:

Section	Details
ML Layer	- Provides statistical propensity score - Based on historical data patterns - Does not incorporate contextual signals
LLM Layer	- Provides contextual, signal-based reasoning - Assigns priority bucket with human-readable rationale - Does not recalculate ML scores
Conflict Handling	Conflicts between ML and LLM are surfaced in conflict_flag

- Conflicts are not auto-resolved — human review required
- Both scores are preserved for transparency

10. Human-in-the-Loop

The system must support human oversight and intervention at the output stage:

- Review of agent-generated prioritisation before actioning
- Manual override capability for priority bucket assignments
- Feedback capture to support future model improvement
- Conflict resolution guided by human judgment

11. Constraints

The following constraints are mandatory and must be enforced across all agents:

Constraint	Description
No ML recalculation in LLM	The Reasoning Agent must not recalculate or override ML scores
No unsupported fields	Agents must not generate or expose fields not defined in the output schema
Structured outputs only	All agent outputs must conform to the defined JSON schema
Conflict visibility	Conflicts must always be surfaced — never silently suppressed

12. Success Criteria

The POC will be considered successful when all of the following criteria are met:

- End-to-end agent flow executes without errors for a representative dataset
- Account prioritisation is meaningful and aligns with business expectations
- Conflicts between ML and LLM outputs are visible in output payload
- Output is explainable and auditable by sales teams
- Human override and feedback mechanisms are functional

13. Future Enhancements

The following capabilities are out of scope for the initial POC but planned for future iterations:

Enhancement	Description
Autonomous Execution	Agents trigger actions without human initiation
CRM Integration	Direct read/write to CRM systems (e.g. Salesforce)
Continuous Learning	Model retraining on feedback and outcome data
Multi-Agent Orchestration	Cross-workflow agent collaboration at enterprise scale

— End of Document —